

Titanic Survival Prediction Project

Project Overview: This Jupyter Notebook documents the Titanic Survival Prediction project, which aims to predict passenger survival during the tragic sinking of the Titanic. We analyze various factors such as socio-economic status, age, gender, and more to develop an accurate survival prediction model.

Project Highlights:

- Data Loading and Exploration
- Data Preprocessing
- Model Training and Evaluation
- Experimenting with Multiple Algorithms
- Selecting and Fine-Tuning the Best Model

Project Objectives:

- Analyze and visualize Titanic dataset
- Build predictive models to determine survival
- Evaluate and compare model performance
- Provide insights for passenger survival factors

Contact Information: LinkedIn: [Asad Ali Mulla's LinkedIn Profile](https://www.linkedin.com/in/asad-ali-mulla-044240262/)
(<https://www.linkedin.com/in/asad-ali-mulla-044240262/>) GitHub: [Asad Ali Mulla's GitHub Profile](https://github.com/Asadxio)
(<https://github.com/Asadxio>)

Task: The task is to create a system that predicts whether a person would survive the Titanic sinking based on factors such as socio-economic status, age, gender, and more.

Data Loading and Exploration

```
In [1]: import pandas as pd
from sklearn.impute import SimpleImputer
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import warnings
warnings.filterwarnings("ignore")
```

```
In [2]: train_data = pd.read_csv('train.csv')
test_data = pd.read_csv('test.csv')
```

```
In [3]: import seaborn as sns
import matplotlib.pyplot as plt

train_data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column        Non-Null Count  Dtype
---  -
 0   PassengerId    891 non-null    int64
 1   Survived       891 non-null    int64
 2   Pclass        891 non-null    int64
 3   Name          891 non-null    object
 4   Sex           891 non-null    object
 5   Age           714 non-null    float64
```

```
In [3]: import seaborn as sns
import matplotlib.pyplot as plt
```

```
train_data.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   PassengerId     891 non-null    int64  
 1   Survived        891 non-null    int64  
 2   Pclass         891 non-null    int64  
 3   Name           891 non-null    object  
 4   Sex            891 non-null    object  
 5   Age            714 non-null    float64 
 6   SibSp         891 non-null    int64  
 7   Parch         891 non-null    int64  
 8   Ticket         891 non-null    object  
 9   Fare          891 non-null    float64 
10   Cabin         204 non-null    object  
11   Embarked      889 non-null    object  
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
In [4]: train_data.head()
```


Out[4]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Ci
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	I
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	I
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	I

```
In [5]: sns.countplot(x='Survived', data=train_data)
plt.title('Survival Distribution (0 = No, 1 = Yes)')
plt.show()
```



```
In [5]: sns.countplot(x='Survived', data=train_data)
plt.title('Survival Distribution (0 = No, 1 = Yes)')
plt.show()
```



```
In [6]: import pandas as pd
from sklearn.impute import SimpleImputer

train_data = pd.read_csv('train.csv')
test_data = pd.read_csv('test.csv')
```

```
In [7]: train_data.drop(['Name', 'Ticket', 'Cabin'], axis=1, inplace=True)
test_data.drop(['Name', 'Ticket', 'Cabin'], axis=1, inplace=True)
```

```
In [8]: num_imputer = SimpleImputer(strategy='mean')
train_data[['Age', 'Fare']] = num_imputer.fit_transform(train_data[['Age', 'Fare']])
test_data[['Age', 'Fare']] = num_imputer.transform(test_data[['Age', 'Fare']])
```

```
In [9]: train_data = pd.get_dummies(train_data, columns=['Sex', 'Embarked'], drop_first=True)
test_data = pd.get_dummies(test_data, columns=['Sex', 'Embarked'], drop_first=True)
```

```
In [10]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Split the data into features (X_train) and the target variable (y_train)
y_train = train_data['Survived']
X_train = train_data.drop('Survived', axis=1)
```

```
In [11]: # Split the training data into training and validation sets (80% training, 20% validation)
X_train, X_valid, y_train, y_valid = train_test_split(X_train, y_train, test_size=0.2, random_state=0)
```

```
In [12]: # Calculate the accuracy on the validation data
model = LogisticRegression()
model.fit(X_train, y_train)
y_valid_pred = model.predict(X_valid)
print(f"Accuracy on Validation Data: {accuracy_score(y_valid, y_valid_pred):.2f}")
```

```
In [13]: # Train the model on the training data
model.fit(X_train, y_train)
```

```
Out[18]: # Print a Classification report for more detailed evaluation
print("\nClassification Report:")
print(classification_report(y_valid, y_valid_pred))
```

```
In [14]: # Make predictions on the validation data
y_valid_pred = model.predict(X_valid)

precision    recall  f1-score   support
```

```
0          0.79          0.84          0.81         105
1          0.75          0.60          0.67          74
```

```
In [12]: # Calculate accuracy with Regressor on data
model = LogisticRegression()
print(f"Accuracy on Validation Data: {accuracy:.2f}")
```

```
In [13]: # Train the model on the training data
model.fit(X_train, y_train)
```

```
Out[16]: # Print a classification report for more detailed evaluation
print("\nClassification Report:")
print(classification_report(y_valid, y_valid_pred))
```

```
In [14]: # Make predictions on the validation data
y_valid_pred = model.predict(X_valid)
```

	precision	recall	f1-score	support
0	0.79	0.84	0.81	105
1	0.75	0.69	0.72	74
accuracy			0.78	179
macro avg	0.77	0.76	0.77	179
weighted avg	0.78	0.78	0.77	179

```
In [17]: # Display a confusion matrix to visualize true positives, true negatives, false
print("\nConfusion Matrix:")
print(confusion_matrix(y_valid, y_valid_pred))
```

Confusion Matrix:

```
[[88 17]
 [23 51]]
```

```
In [18]: # Initialize different classification models
models = [
    ('Logistic Regression', LogisticRegression()),
    ('Random Forest', RandomForestClassifier(random_state=42)),
    ('Support Vector Machine', SVC()),
    ('K-Nearest Neighbors', KNeighborsClassifier()),
    ('Decision Tree', DecisionTreeClassifier(random_state=42))
]
```

```
In [19]: # Evaluate each model using cross-validation
for name, model in models:
    scores = cross_val_score(model, X_train, y_train, cv=5, scoring='accuracy')
    mean_accuracy = scores.mean()
    std_accuracy = scores.std()

    print(f"{name}:")
    print(f"Mean Accuracy: {mean_accuracy:.2f}")
    print(f"Standard Deviation: {std_accuracy:.2f}")
    print()
```

Logistic Regression:
Mean Accuracy: 0.77
Standard Deviation: 0.04

Random Forest:
Mean Accuracy: 0.81
Standard Deviation: 0.01

Support Vector Machine:

```
In [21]: Mean Accuracy: 0.65
best_model = RandomForestClassifier(random_state=42)
std_deviation: 0.01
best_model.fit(X_train, y_train)
y_valid_pred = best_model.predict(X_valid)
K-Nearest Neighbors:
Mean Accuracy: 0.62
```

```
In [22]: Standard Deviation: 0.03
print(f"Accuracy on Validation Data: {accuracy:.2f}")
print("\nClassification Report:")
print(classification_report(y_valid, y_valid_pred))
print(f"Standard Deviation: {std_deviation:.2f}")
print(confusion_matrix(y_valid, y_valid_pred))
```

Accuracy on Validation Data: 0.83

Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

```
In [21]: Mean Accuracy: 0.65
best_model = RandomForestClassifier(random_state=42)
Standard Deviation: 0.01
best_model.fit(X_train, y_train)
y_valid_pred = best_model.predict(X_valid)
K-Nearest Neighbors:

In [22]: Mean Accuracy: 0.62
Standard Deviation: 0.003e(y_valid, y_valid_pred)
print(f"Accuracy on Validation Data: {accuracy:.2f}")
print(f"Confusion Matrix Report:")
print(f"Classification Report:")
print(f"Mean Accuracy: {accuracy:.2f}")
print(f"Standard Deviation: {std:.2f}")
print(confusion_matrix(y_valid, y_valid_pred))
```

Accuracy on Validation Data: 0.83

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.90	0.86	105
1	0.83	0.74	0.79	74
accuracy			0.83	179
macro avg	0.83	0.82	0.82	179
weighted avg	0.83	0.83	0.83	179

Confusion Matrix:
[[94 11]
 [19 55]]